LLM Course

LLM Course documentation

Ungraded quiz

LLM Course

■ View all resources

Agents Course

Audio Course

Community Computer Vision Course

Deep RL Course

Diffusion Course

LLM Course

MCP Course

ML for 3D Course

ML for Games Course

Open-Source AI Cookbook

Robotics Course

a smol course

Search documentation

AR

BN

DE

EN

ES

FA

FR

GJ

HE

HI

ID

IT

JA

KO

MY

NE

PL

PT

RO

RU

RUM

TE

TH

TR

VI

ZH-CN

ZH-TW

Join the Hugging Face community

and get access to the augmented documentation experience

Collaborate on models, datasets and Spaces

Faster examples with accelerated inference

Switch between documentation themes

Sign Up

to get started

Copy page

Ungraded quiz

So far, this chapter has covered a lot of ground! Don’t worry if you didn’t grasp all the details, but it’s to reflect on what you’ve learned so far with a quiz.

This quiz is ungraded, so you can try it as many times as you want. If you struggle with some questions, follow the tips and revisit the material. You’ll be quizzed on this material again in the certification exam.

1. Explore the Hub and look for the roberta-large-mnli checkpoint. What task does it perform?

Summarization

Text classification

Text generation

Submit

2. What will the following code return?

Copied

from

transformers

import

pipeline

ner = pipeline(

"ner"

, grouped_entities=

True

)

ner(

"My name is Sylvain and I work at Hugging Face in Brooklyn."

)

It will return classification scores for this sentence, with labels "positive" or "negative".

It will return a generated text completing this sentence.

It will return the words representing persons, organizations or locations.

Submit

3. What should replace … in this code sample?

Copied

from

transformers

import

pipeline

filler = pipeline(

"fill-mask"

, model=

"bert-base-cased"

)

result = filler(

"..."

)

This <mask> has been waiting for you.

This [MASK] has been waiting for you.

This man has been waiting for you.

Submit

4. Why will this code fail?

Copied

```
from
transformers
import
pipeline
classifier = pipeline(
"zero-shot-classification"
)
result = classifier(
"This is a course about the Transformers library"
)
```

This pipeline requires that labels be given to classify this text.

This pipeline requires several sentences, not just one.

The ■ Transformers library is broken, as usual.

This pipeline requires longer inputs; this one is too short.

Submit

5. What does “transfer learning” mean?

Transferring the knowledge of a pretrained model to a new model by training it on the same dataset.

Transferring the knowledge of a pretrained model to a new model by initializing the second model with the first model's weights.

Transferring the knowledge of a pretrained model to a new model by building the second model with the same architecture as the first model.

Submit

6. True or false? A language model usually does not need labels for its pretraining.

True

False

Submit

7. Select the sentence that best describes the terms “model”, “architecture”, and “weights”.

If a model is a building, its architecture is the blueprint and the weights are the people living inside.

An architecture is a map to build a model and its weights are the cities represented on the map.

An architecture is a succession of mathematical functions to build a model and its weights are those functions parameters.

Submit

8. Which of these types of models would you use for completing prompts with generated text?

An encoder model

A decoder model

A sequence-to-sequence model

Submit

9. Which of those types of models would you use for summarizing texts?

An encoder model

A decoder model

A sequence-to-sequence model

Submit

10. Which of these types of models would you use for classifying text inputs according to certain labels?

An encoder model

A decoder model

A sequence-to-sequence model

Submit

11. What possible source can the bias observed in a model have?

The model is a fine-tuned version of a pretrained model and it picked up its bias from it.

The data the model was trained on is biased.

The metric the model was optimizing for is biased.

Submit

Update

on GitHub

←

Transformer Architectures

Inference with LLMs

→